


ECE2810J

Data Structures and Algorithms

- ▶ Hashing: Basics and Hash Function
- ▶ Learning Objectives:
 - ▶ Know the purpose of hashing
 - ▶ Understand the collision problem
 - ▶ Know how to design a good hash function



Outline



REVIEW OF
DICTIONARY



HASHING BASICS



HASH FUNCTION

Dictionary

- ▶ How do you use a dictionary?
 - ▶ Look up a “word” and find its meaning.
- ▶ We also have an abstract data type called dictionary.
 - ▶ It is a collection of pairs, each containing a **key** and an **value**
(**key**, **value**)
 - ▶ **Important:** Different pairs have different keys.



Dictionary

- ▶ Key space is usually more regular/structured than value space, so easier to search.
- ▶ Dictionary is optimized to quickly **add** (**key**, **value**) pair and **retrieve value** by key.

Methods

- ▶ **Value find(Key k):** Return the value whose key is **k**. Return **Null** if none.
- ▶ **void insert(Key k, Value v):** Insert a pair (**k**, **v**) into the dictionary. If the pair with key as **k** already exists, update its value.
- ▶ **Value remove(Key k):** Remove the pair with key as **k** from the dictionary and return its value. Return **Null** if none.

Runtime for Array Implementation

Pair Array[MAXSIZE] :

a	b	c	d			
---	---	---	---	--	--	--

- ▶ Unsorted array
 - ▶ find() $O(n)$
 - ▶ insert() $O(n)$: $O(n)$ to verify duplicate, $O(1)$ to put at the end
 - ▶ remove() $O(n)$: $O(n)$ to verify existence, $O(1)$ to exchange the “hole” with the last element
- ▶ Sorted array
 - ▶ find() $O(\log n)$: binary search
 - ▶ insert() $O(n)$: $O(\log n)$ to verify duplicate, $O(n)$ to insert
 - ▶ remove() $O(n)$: $O(\log n)$ to verify existence, $O(n)$ to remove

Can we do **find**, **insert**, and **remove** in $O(1)$ time?

Outline



Review of Dictionary



Hashing Basics



Hash Function


Hashing: High-Level Idea

- ▶ **Setup:** A universe U of objects
 - ▶ E.g., All names, all IP addresses, etc.
 - ▶ Generally, very BIG!
- ▶ **Goal:** Want to maintain an evolving set $S \subseteq U$
 - ▶ E.g., 200 students, 500 IP addresses
 - ▶ Generally, of reasonable size.
- ▶ Naïve solutions
 1. Array-based solution (index by $u \in U$)
 - ▶ $\Theta(1)$ operation time, BUT $\Theta(|U|)$ space.
 2. Linked list-based solution:
 - ▶ $\Theta(|S|)$ space, BUT $\Theta(|S|)$ operation time.

Can we get the best of both solutions?

Hashing: High-Level Idea

- ▶ Solution:
 - ▶ Pick an array A of n buckets.
 - ▶ $n = c|S|$: a small multiple of $|S|$.
 - ▶ Choose a hash function $h: U \rightarrow \{0, 1, \dots, n - 1\}$
 - ▶ h is fast to compute.
 - ▶ The same key is always mapped to the **same** location.
 - ▶ Store item k in $A[h(k)]$
- ▶ The array is called **hash table**
 - ▶ An array of **buckets**, where each bucket contains items as assigned by a hash function.
 - ▶ $h[k]$ is called the **home bucket** of key k .



Hashing Example

- Pairs are: (22,a), (33,b), (3,c), (73,d), (85,e)
- Hash table is $A[0:7]$ and table size is $M = 8$
- Hash function is $h[key] = key/11$
- Every item with key is stored in the bucket $A[h(key)]$

(3,c)		(22,a)	(33,b)			(73,d)	(85,e)
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]

Question: What is the time complexity for `find()`, `insert()`, and `remove()`?

$O(1)$

Achieves both fast runtime and efficient memory usage

What Can Go Wrong?

(3,c)		(22,a)	(33,b)			(73,d)	(85,e)
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]

- ▶ Where does (35, g) go?
- ▶ Problem: The home bucket for (35, g) is already occupied!
 - ▶ This is a “collision”.

Collision and Collision Resolution

- ▶ Collision occurs when the hash function maps two or more items—all having **different** search keys—into the **same** bucket.
- ▶ What to do when there is a collision?
 - ▶ **Collision-resolution scheme**: assigns distinct locations in the hash table to items involved in a collision.
- ▶ Two major schemes:
 - ▶ Separate chaining
 - ▶ Open addressing

Insight of Collision: Birthday Problem

- Consider n people with random birthdays (i.e., with each day of the year equally likely). What is the smallest n so that there is at least a 50% chance that two people have the same birthday?
 - A. 23
 - B. 57
 - C. 184
 - D. 367

Collision is inevitable!

Insight of Collision: Birthday Problem

- Consider n people with random birthdays (i.e., with each day of the year equally likely). What is the smallest n so that there is at least a 50% chance that two people have the same birthday?

- A. 23
- B. 57
- C. 184
- D. 367

We want to find the smallest n for which

$$1 - \prod_{i=0}^{n-1} \frac{365-i}{365} > 0.5$$

Collision is inevitable!

Hash Table Issues

- ▶ Choice of the hash function.
- ▶ Collision resolution scheme.
- ▶ Size of the hash table and rehashing.

Outline



REVIEW OF
DICTIONARY



HASHING BASICS



HASH FUNCTION

Hash Function Design Criteria

- ▶ Must compute a bucket for every key in the universe.
- ▶ Must compute the same bucket for the same key.
- ▶ Should be easy and quick to compute.
- ▶ Minimizes collision

The hardest criterion



- ▶ Spread keys out evenly in hash table

- ▶ **Gold standard:** completely random hashing

- ▶ The probability that a randomly selected key has bucket i as its home bucket is $1/n$, $0 \leq i < n$.
 - ▶ Completely random hashing **minimizes** the likelihood of a collision when keys are selected at random.
 - ▶ However, completely random hashing is infeasible due to the need to remember the random bucket.

Bad Hash Functions

- ▶ Example: keys = phone number in China (11 digits)
 - ▶ $|U| = 10^{11}$
 - ▶ **Terrible** hash function: $h(key)$ = first 3 digits of key , i.e., area code
 - ▶ The keys are not spread out evenly. Buckets 010, 021 may have a lot of keys mapped to them, while some buckets have no keys.
 - ▶ **Mediocre** hash function: $h(key)$ = last 3 digits of key .
 - ▶ Still vulnerable to patterns in last 3 digits.

Hash Functions

- ▶ Hash function ($h(key)$) maps key to buckets in two steps:
 1. Convert key into an integer in case the key is not an integer.
 - ▶ A function $t(key)$ which returns an integer value, known as **hash code**.
 2. **Compression map**: Map an integer (hash code) into a home bucket.
 - ▶ A function $c(hashcode)$ which gives an integer in the range $[0, n - 1]$, where n is the number of buckets in the table.
- ▶ In summary, $h(key) = c(t(key))$, which gives an index in the table.

Map Non-integers into Hash Code

- ▶ String: use the ASCII (or UTF-8) encoding of each char and then perform arithmetic on them.
- ▶ Floating-point number: treat it as a string of bits.
- ▶ Images, (viral) code snippets, (malicious) Web site URLs: in general, treat the representation as a bit-string, using all of it or **extracting** parts of it (i.e., www.sjtu.edu.cn).

Strings to Integers

- ▶ Simple scheme: adds up all the ASCII codes for all the chars in the string.
 - ▶ Example: $t(\text{"He"}) = 72 + 101 = 173$.
- ▶ Not good. Why?
 - ▶ Consider English words “post”, “pots”, “spot”, “stop”, “tops”.

Strings to Integers

- A better strategy: Polynomial hash code taking **positional** info into account.

$$t(s[]) = s[0]a^{k-1} + s[1]a^{k-2} + \dots + s[k-2]a + s[k-1]$$

where a is a constant.

- If $a = 33$, the hash codes for “post” and “stop” are
 $t(\text{post}) = 112 \cdot 33^3 + 111 \cdot 33^2 + 115 \cdot 33 + 116 = 4149734$
 $t(\text{stop}) = 115 \cdot 33^3 + 116 \cdot 33^2 + 111 \cdot 33 + 112 = 4262854$

Strings to Integers

$$t(s[]) = s[0]a^{k-1} + s[1]a^{k-2} + \dots + s[k-2]a + s[k-1]$$

- ▶ Good choice of a for English words: 31, 33, 37, 39, 41
 - ▶ What does it mean for a to be a **good** choice? Why are these particular values **good**?
 - ▶ Answer: according to statistics on 50,000 English words, each of these constants will produce less than 7 collisions.
- ▶ In Java, its `String` class has a built-in `hashCode()` function. It takes $a = 31$. Why?
 - ▶ Multiplication by 31 can be replaced by a shift and a subtraction for **better performance**:

`31*i == (i << 5) - i`

Hash function criteria: Should be easy and quick to compute.

Compression Map

- ▶ Map an integer (hash code) into a home bucket.
- ▶ The most common method is by **modulo arithmetic**.

`homeBucket = c(hashcode) = hashcode % n`
where n is the **number of buckets** in the hash table.

- ▶ Example: Pairs are (22,a), (33,b), (3,c), (55,d), (79,e). Hash table size is 7.

	(22,a)	(79,e)	(3,c)		(33,b)	(55,d)
[0]	[1]	[2]	[3]	[4]	[5]	[6]

Hashing by Modulo

- ▶ In practice, keys of an application tend to have a specific pattern.
 - ▶ For example, memory address in computer is multiple of 4.
- ▶ The choice of the hash table size n will affect the distribution of home buckets.

Hashing by Modulo

- ▶ Suppose the keys of an application are more likely to be mapped into even integers.
 - ▶ E.g., memory address is always a multiple of 4.
- ▶ When the hash table size n is an **even** number, **even** integers are hashed into **even** home buckets.
 - ▶ E.g., $n = 14$: $20\%14 = 6$, $32\%14 = 4$, $8\%14 = 8$
- ▶ The bias in the keys results in a bias toward the **even** home buckets.
 - ▶ All **odd** buckets are guaranteed to be empty.
 - ▶ The distribution of home buckets is not uniform!

Hashing by Modulo

- ▶ However, when the hash table size n is **odd**, even (or odd) integers may be hashed into both odd and even home buckets.
 - ▶ E.g., $n = 15$: $20\%15 = 5$, $32\%15 = 2$, $8\%15 = 8$
 $15\%15 = 0$, $3\%15 = 3$, $23\%15 = 8$
- ▶ The bias in the keys does not result in a bias toward either the odd or even home buckets.
 - ▶ Better chance of uniform distribution of home buckets.
- ▶ So do not use an even hash table size n .

Hashing by Modulo

- ▶ Similar **biased** distribution of home buckets happens in practice when the hash table size n is a multiple of small prime numbers.
- ▶ The effect of each prime divisor p of n **decreases** as p gets **larger**.
- ▶ Ideally, choose the hash table size n as a **large prime number**.

Code Exercise: Hashing by Modulo (~15 mins)

- Code can be found in :
 - Canvas -> Code Exercise -> Hashing